

# Elasticsearch Plugin - Reference Documentation

Carsten Götzinger

Version 2.7.4

# Table of Contents

1. Introduction	1
1.1. Features	1
1.2. History	1
1.3. Acknowledgments	6
1.4. Plugin Versioning	7
1.5. License	8
2. Configuration	9
2.1. Overriding Spring Boot elasticsearch version	9
2.2. Default configuration script	9
2.3. Client mode	12
2.4. Mapping Migration properties	14
2.5. Dynamic Method Injection	15
2.6. Others properties	16
3. Quick Start	20
3.1. Default mapping	21
3.2. Custom mapping	22
3.3. Limit properties with only/except	22
3.4. Including transients	23
3.5. Transients and collections	24
3.6. The Mapper Attachment plugin	25
4. Mapping	26
4.1. Class Mapping	26
4.2. Properties Mapping	27
4.3. Parent Child	29
4.4. Geo Point	30
4.5. Alias	30
4.6. Dynamic	30
4.7. Searchable Component Reference	31
4.8. Mapping Migrations	34
5. Indexing	38
5.1. Index examples	38
5.2. Unindex examples	38
6. Searching	40
6.1. Query Strings	41
6.2. Query Closure	41
6.3. Query Builder	41
6.4. Filter Closure	42
6.5. Filter Builder	42

6.6. Highlighting .....	42
6.7. Sorting .....	44
6.8. Pagination .....	45
6.9. Aggregation .....	45
7. Admin .....	48
7.1. Refresh .....	48
7.2. Delete Index .....	48
8. Low Level API .....	49
9. Example .....	50
9.1. Twitter .....	50
9.2. Geo Distance Search .....	53
9.3. Parent/Child mapping .....	55
10. Frequently Asked Questions .....	56
10.1. Conflicting module versions .....	56

# Chapter 1. Introduction

The ElasticSearch plugin intends to implement a simple integration with Grails of the Open Source Search Engine [ElasticSearch](#), which is based on Lucene and provide distributed capabilities.

The plugin focuses on exposing Grails domain classes for the moment. It highly takes the existing [Searchable Plugin](#) as reference for its syntax and behaviour.

Note that the plugin is still under development, so you may not be able to use all the features of ElasticSearch yet.

In addition to this document, you may want to read the official ElasticSearch documentation [here](#).

## 1.1. Features

- Maps domain classes to their corresponding index in ElasticSearch
- Provides an ElasticSearch service for cross-domain searching
- Injects domain class methods for specific domain searching, indexing and unindexing
- Automatically mirrors any changes made through GORM to the index
- Allow to use the Groovy Content Builder DSL for search queries
- Support for term highlighting

## 1.2. History

### 1.2.1. Grails 4.x version

- December 24, 2019
  - 2.8.0
    - Upgraded to Grails 4.0.1
    - Upgraded to Elasticsearch 7.5.1

### 1.2.2. Grails 3.x version

- February 12, 2020
  - 2.7.3
    - Ignore `_domainTypeName` in `DomainClassUnmarshaller`
- January 2, 2020
  - 2.7.2
    - Upgraded to Grails 3.3.11
    - Upgraded to Elasticsearch 7.5.1
    - Don't use old `_type` in mapping (deprecated in ES)

- Optionally add `_domainTypeName` in mapping
- August 12, 2019
  - 2.7.0
    - Upgraded to Grails 3.3.10
    - Upgraded to Elasticsearch 7.3.0
      - Return type of count method in `ElasticSearchService` is now `java.lang.Long` (was `java.lang.Integer` before)
      - The property `total` in class `grails.plugins.elasticsearch.ElasticSearchResult` is now of type `org.apache.lucene.search.TotalHits`
    - Make java.time classes usable
    - Make aggregations usable with this plugin
- August 6, 2019
  - 2.6.0
    - Upgraded to Grails 3.3.9
    - Upgraded to Elasticsearch 6.8.2
      - mapping option `index`: was before `analyzed`, `not_analyzed`, `no` and can now only be: `true`, `false`
      - `include_in_all` is deprecated because of <https://www.elastic.co/guide/en/elasticsearch/reference/6.4/mapping-all-field.html>, but could be enabled with <https://www.elastic.co/guide/en/elasticsearch/reference/6.4/mapping-all-field.html#enabling-all-field>, but now completely disabled in `ElasticSearchMappingFactory.groovy:125`
      - Removal of mapping types: <https://www.elastic.co/guide/en/elasticsearch/reference/current/removal-of-types.html>
        - In the documentation 4.1. Class Mapping → 4.1.1. `root` is not used anymore, because every domain class will get its own index
      - `parent` is no longer useable (not supported by Elasticsearch 6.x)
      - Embedded ES was removed and so there is no accurate possibility to run integration tests with ES, see here: <https://discuss.elastic.co/t/in-memory-testing-with-resthighlevelclient/106196/6> . For now the integration tests are only working if a local elasticsearch version is running, which can be downloaded here: <https://www.elastic.co/downloads/past-releases> - please note that you've to use the same version like described in the above comments, so for now it is 6.8.2
  - May 10, 2018
    - 2.4.1
      - `search` method now returns `ElasticsearchResult` instead of `Map`.
      - Added configuration `elasticSearch.plugin.mapperAttachment.enabled` to disable Mapper Attachment plugin.
  - December 12, 2017

- 2.4.0.RC1
  - Support Grails 3.3.x
  - The mapper-attachments plugin has been changed to **provided**.
- August 17, 2017
  - 1.4.1
    - Upgraded to Elasticsearch 5.4.1
      - Changed data types 'string' to 'text' and 'keyword' with sensible defaults
      - BREAKING CHANGE: if you want to sort or aggregate by text properties you need to add `fielddata: true` since it is disabled by default in ES 5.x
      - BREAKING CHANGE: default search type changed to `QUERY_THEN_FETCH`. See the tip at the end of [here](#)
      - Added initial support for [fields](#)
      - Ignore `_all` configuration for component objects
- June 3, 2016
  - 1.2.0
    - Upgraded to Elasticsearch 2.3.3
      - Support for **elasticsearch-groovy** client removed (no available compatible version, pending on [Issue 35](#))
      - **delete** migration strategy has been replaced by **deleteIndex**, due to new Elasticsearch 2.x API's
    - Removed test classes from plugin distribution
- April 18, 2016
  - 1.0.0.1
    - Add ability to change search method name in domain class via config
    - Updated documentation to asciidoc
- February 29, 2016
  - 1.0.0
    - Support for Grails 3.1.1

### 1.2.3. Grails 2.x version

- April 3, 2016
  - 0.1.0
    - New Elasticsearch 2.1.2 release
    - Immutable Settings Removed - Use **Settings.builder()** instead of **ImmutableSettings.builder()**
    - **BroadcastOperationResponse** got renamed to **BroadcastResponse**

- Removed `deleteMapping`
- `memory` type is now deprecated
- Query/filter refactoring - `org.elasticsearch.index.queries.FilterBuilders` has been removed as part of the merge of queries and filters. These filters are now available in `QueryBuilders` with the same name. All methods that used to accept a `FilterBuilder` now accept a `QueryBuilder` instead.
- For more information related to changes in underline Java API [click here](#)
- June 30, 2015
  - 0.0.4.5
    - Upgrade to Elasticsearch 1.6.0
    - Support the return of aggregation results
- June 15, 2015
  - 0.0.4.5
    - Add the ability to define property names that are excluded by default
    - Fix NPE
    - Add the attachment type
- March 5, 2015
  - 0.0.4.4
    - Upgrade to Elasticsearch-Groovy 1.4.4
- February 22, 2015
  - 0.0.4.3
    - Add mapping configuration support for '\_all'
    - Fix issue with indexing nested GeoPoint
    - Add support for transient properties
- February 10, 2015
  - 0.0.4.2
    - Reduce severity of non-searchable property in index document when unmarshalling domain
- February 03, 2015
  - 0.0.4.1
    - Upgrade to Elasticsearch 1.4.2
    - Enable configuration of the number of replicas created per shard
- January 28, 2015
  - 0.0.4.0
    - Included Mapping migrations
    - Included read and write aliases to indices to deal with migrations on multinode

## deployments

- December 14, 2014
  - 0.0.3.8
    - Upgrade to ElasticSearch 1.4.1
    - Support the min\_score query parameter.
    - Try to detect the MongoDB without using the plugin manager.
- December 01, 2014
  - 0.0.3.7
    - Create separate SimpleTypeConverter per-thread
- November 06, 2014
  - 0.0.3.6
    - Upgrade to ElasticSearch 1.4.0
- October 28, 2014
  - 0.0.3.5
    - Fix the bulk index query iteration.
- October 14, 2014
  - 0.0.3.4
    - Upgrade to latest version of ElasticSearch and remove the Groovy client dependency.
- August 28, 2014
  - 0.0.3.3
    - Configure a component field to act as an inner object instead of a nested object.
- August 3, 2014
  - 0.0.3.2
    - Add the ability to mark fields with aliases
    - Support ES client HTTP configuration parameters
    - Improve Hibernate 4 support
- June 9, 2014
  - 0.0.3.1
    - Upgrade to ElasticSearch 1.2.x
    - Add special treatment for MongoDB ObjectId data types
    - Return raw result objects when no class mapping is found
    - Fix integration-test NPE
- May 25, 2014
  - 0.0.3.0
    - Upgrade to Grails dependency 2.2.x



- Upgrade to Grails runtime 2.3.x
- Upgrade to Elasticsearch 1.x
- Apply Elasticsearch 1.x compatibility fixes
- Enable customization of index name types when mapping classes
- May 15, 2014
  - 0.0.2.6
    - Use 'grails.util.Holders' instead of ApplicationHolder
- April 2, 2014
  - 0.0.2.5
    - Start releasing the plugin as 'elasticsearch' instead of 'elasticsearch-gorm'
    - Fix NPE when marshalling JSONObject fields
- March 24, 2014
  - 0.0.2.4
    - GeoPoint mapping
    - Injected service now supports filters (e.g. geo\_reference) and sort builders (e.g. for geo\_distance sorting)
    - Marshalled date values are now with correct time zone
    - Removed dependency on Java 7
    - Fix support of BigDecimal
    - Searchable mapping property name and Elasticsearch plugin path are now configurable.
- February 4, 2014
  - 0.0.2.3 Bugfix release
- January 19, 2014
  - 0.0.2.2 Bugfix release
- November 24, 2013
  - 0.0.2.1 Bugfix release
- November 12, 2013
  - 0.0.2 release
- November 2, 2013
  - initial 0.0.1 release

## 1.3. Acknowledgments

Many thanks to all the users who reported issues and sent me pull requests.

### 1.3.1. Authors and Contributors

- [Noam Y. Tenne](#)
- [Stefan Rother-Stübs \(Dating Cafe\)](#)
- [Sven Kiesewetter \(Dating Cafe\)](#)
- Michael Schwartz (Dating Cafe)
- [Marcos Carceles](#)
- [Puneet Behl](#)
- [Carsten Götzinger \(cgo IT\)](#)

### 1.3.2. Authors and Contributors of the original plugin

- [Manuarii Stein \(doc4web consulting\)](#)
- [Stephane Maldini \(doc4web consulting\)](#)
- [Serge P. Nekoval](#)

### 1.3.3. Previous work

Graeme Rocher started the first draft which this plugin is based on.

Get the full and updated list of contributors on the [github](#) repository.

## 1.4. Plugin Versioning



The versioning model has changed. The version number of the plugin will reflect the one of the underline integrated Elasticsearch. If necessary a 4th level point release number will be used for successive changes on the plugin's code with same version of Elasticsearch.

`<GRAILS_VERSION>.<ES_VERSION>.<FEATURE/PATCH_VERSION>`, where there isn't really a 1-to-1 plugin version to grails or es version, but we just increase our major or minor version by one, whenever there are breaking changes on either Grails or ES. `<GRAILS_VERSION>.<ES_VERSION>.<FEATURE/PATCH_VERSION>`, where there isn't really a 1-to-1 plugin version to grails or es version, but we just increase our major or minor version by one, whenever there are breaking changes on either Grails or ES. Therefore future release versions could look something similar to this (depending on Grails and Elasticsearch versions):

Plugin Version	Grails	Elasticsearch
2.8.0	4.0.x	7.x
2.7.3	3.3.x	7.x
2.7.2	3.3.x	7.x
2.7.0	3.3.x	7.x
2.6.0	3.3.x	6.x

2.4.0	3.3.x	5.x
1.4.1	3.1.x	5.x (tested with 5.4.1)
1.2.0	3.1.x	2.3.x
1.0.0.2	3.0.x	1.x
0.1.0	2.4.x	2.1.x
0.0.4.5	2.4.x	1.x
0.2.x	2.4.x	2.3.x
0.3.x	2.4.x	(hypothetical) 2.4.x
0.5.x	2.4.x	(hypothetical) 5.0.x
1.3.x	3.1.x	(hypothetical) 2.4.x
1.5.x	3.1.x	(hypothetical) 5.0.x
2.3.x	(hypothetical) 3.2.x	(hypothetical) 2.3.x
3.3.x	(hypothetical) 4.0.x	(hypothetical) 2.3.x

*Existing versions in **bold***

The current version is **2.7.4** (for Grails 2.x the latest version is **0.1.0**)

## 1.5. License

This plugin is released under the [Apache License, Version 2.0](#)

# Chapter 2. Configuration

The plugin provide a default configuration, but you may add your own settings in your **Config.groovy** for Grails 2.x and **application.groovy** or **application.yml** for Grails 3.x.

## 2.1. Overriding Spring Boot elasticsearch version

Spring Boot 1.3.x supports Elasticsearch 1.5.2 OOTB (<https://github.com/spring-projects/spring-boot/blob/master/spring-boot-dependencies/pom.xml#L76>) and will install dependencies for this version, if not explicitly overridden. To do so add the following to your **build.gradle**:

```
def elasticsearchVersion = '5.4.3'
ext['elasticsearch.version'] = elasticsearchVersion
```

## 2.2. Default configuration script

Below is the default configuration loaded by the plugin (any of your settings in the Config.groovy script overwrite those).

*Config.groovy*

```
elasticsearch {
    /**
     * Date formats used by the unmarshaller of the JSON responses
     */
    date.formats = ["yyyy-MM-dd'T'HH:mm:ss'Z'"]

    /**
     * Hosts for remote ElasticSearch instances.
     * Will only be used with the "transport" client mode.
     * If the client mode is set to "transport" and no hosts are defined, ["localhost",
     9300] will be used by default.
     */
    client.hosts = [
        [host:'localhost', port:9300]
    ]

    /**
     * Default mapping property exclusions
     *
     * No properties matching the given names will be mapped by default
     * i.e., when using "searchable = true"
     *
     * This does not apply for classes using mapping by closure
     */
    defaultExcludedProperties = ["password"]
}
```

```

    * Determines if the plugin should reflect any database save/update/delete
    automatically
    * on the ES instance. Default to false.
    */
    disableAutoIndex = false

/**
 * Should the database be indexed at startup.
 *
 * The value may be a boolean true|false.
 * Indexing is always asynchronous (compared to Searchable plugin) and executed
    after BootStrap.groovy.
    */
    bulkIndexOnStartup = true

/**
 * Max number of requests to process at once. Reduce this value if you have memory
    issue when indexing a big amount of data
 * at once. If this setting is not specified, 500 will be use by default.
    */
    maxBulkRequest = 500

/**
 * The name of the Elasticsearch mapping configuration property that annotates
    domain classes. The default is 'searchable'.
    */
    searchableProperty.name = 'searchable'
}

environments {
    development {
        /**
         * Possible values : "local", "dataNode", "transport"
         *
         */
        elasticSearch.client.mode = 'local'
    }
    test {
        elasticSearch {
            client.mode = 'local'
            index.store.type = 'memory' // store local node in memory and not on disk
        }
    }
    production {
        elasticSearch.client.mode = 'transport'
    }
}

```

## 2.2.1. Grails 3.3.x Sample YAML configuration

*application.yml*

```
elasticSearch:
  date:
    formats: ["yyyy-MM-dd'T'HH:mm:ss.SSS'Z'"]
  client.hosts:
    - {host: localhost, port: 9300}
  defaultExcludedProperties: ['password']
  disableAutoIndex: false
  index:
    compound_format: true
  unmarshallComponents: true
  searchableProperty:
    name: searchable
  includeTransients: false
environments:
  development:
    elasticSearch:
      client:
        mode: local
        transport.sniff: true
        bulkIndexOnStartup: true
  test:
    elasticSearch:
      client:
        mode: local
        transport.sniff: true
      dataStoreImpl: hibernateDatastore
      index:
        store.type: memory
        settings:
          analysis:
            filter:
              replace_synonyms:
                type: synonym
                synonyms: ['abc => xyz']
            analyzer:
              test_analyzer:
                tokenizer: standard
                filter: ['lowercase']
              repl_analyzer:
                tokenizer: standard
                filter: ['lowercase', 'replace_synonyms']

  production:
    elasticSearch:
      client:
        mode: transport
```

## 2.3. Client mode

You can set the plugin in 3 different modes, detailed on the [official ElasticSearch doc](#). The mode is defined with the following config key:

*applicaiton.groovy or Config.groovy*

```
elasticSearch.client.mode = '<mode>'
```

*application.yml*

```
elasticSearch:
  client:
    mode: <mode>
```

Table 1. Possible values for client node

Value	Description
node (Removed in version 2.4.1 of the plugin)	The plugin create its own node and join the ElasticSearch cluster as a client node ( <code>node.client = true</code> ). This setting requires that you have an ElasticSearch instance running and available on your network (use the discovery feature)
dataNode	The plugin create its own node and join the ElasticSearch cluster as a node that can hold data. This setting requires that you have an ElasticSearch instance running and available on your network (use the discovery feature)
local	The plugin create its own local (to the JVM) node. Does not require any running ElasticSearch instance. Useful for development or testing.
transport	The plugin create a transport client that will connect to a remote ElasticSearch instance without joining the cluster.

"Transport" mode needs you to provide the host address and port. You can define one or multiple hosts with the following config key:

*application.groovy or Config.groovy*

```
elasticSearch.client.hosts = [
    [host:'192.168.0.3', port:9300],
    [host:'228.168.0.4', port:9300]
]
```

*application.yml*

```
elasticsearch:
  client:
    hosts:
      - {host: 192.168.0.3, port: 9300}
      - {host: 228.168.0.4, port: 9300}
```

If no host is defined, `localhost:9300` will be used by the transport client.

### 2.3.1. Using node or dataNode during development

The `node.client` setting has been removed. A node with such a setting set will not start up. Instead, each node role needs to be set separately using the existing `node.master`, `node.data` and `node.ingest` supported static settings.

If you configure your dev environment to use `node` or `dataNode` you might see the following exception:

```
'elasticsearchClient': FactoryBean threw exception on object creation; nested
exception is java.lang.IllegalStateException: path.home is not configured
```

In order to make this work you need to define `es.path.home` in VM options.

*Grails 2.x*

```
grails run-app -Des.path.home=<PATH_TO_ELASTICSEARCH_HOME_DIR>
```

*Grails 3.x (build.gradle)*

```
bootRun {
    jvmArgs = ['-Des.path.home=<PATH_TO_ELASTICSEARCH_HOME_DIR>']
}
```

### 2.3.2. Disable Mapper Attachment plugin

The Mapper Attachment plugin is Deprecated since version 5.0.0 and will be replaced by [ingest-attachment](#) plugin.

So, if you are using ES plugin with mode: `local` and do not want to use Mapper Attachment plugin then you can disable the plugin by using following configuration:



```

elasticsearch:
  plugin:
    mapperAttachment:
      enabled: false

```



Please note that by default the plugin is enabled.

## 2.4. Mapping Migration properties

Define the application's behaviour when a conflict is found while installing Elasticsearch mappings on startup. For a detailed explanation, see [Mapping Migrations](#).

### 2.4.1. `elasticsearch.migration.strategy`

Defines the behaviour to follow if an error occurs on startup when the application is installing new mappings on Elasticsearch due to conflicting mappings.

Table 2. Possible Values for migration strategy

Value	Description
'none'	No changes on the indices or mappings will happen, the merge problem will be logged and a <code>MappingException</code> will be thrown.
'delete'	The conflicting mapping will be deleted (along with all indexed content of that type) and replaced with a new mapping. Deleted content can be automatically reindexed on startup by using this in combination the <code>elasticsearch.bulkIndexOnStartup</code> config option
'alias'	Applies <a href="#">Elasticsearch recommended approach for migrating conflicting mappings</a> . A new numbered index will be created ( <code>&lt;indexName&gt;_vX</code> ) where new mappings will be installed for all the types included on the original index. An Elasticsearch alias called <code>&lt;indexName&gt;</code> will point to the new index. As content won't be available on the new index, content can be automatically reindexed on startup by using this in combination the <code>elasticsearch.bulkIndexOnStartup</code> config option. It is recommended to set <code>elasticsearch.aliasReplacesIndex</code> to deal with potential index/alias conflicts.



The default is 'alias'.

## 2.4.2. `elasticSearch.migration.aliasReplacesIndex`

Deals with a special conflict case using the 'alias' strategy. When the 'alias' migration strategy is chosen and there's a mapping conflict on an index, defines whether to replace the index with a versioned index (`<indexName>_vX`) and an alias (`<indexName>`). This is required when applying the alias strategy on top of existing indices for the first time as indices cannot be renamed (from `<indexName>` to `<indexName>_vX`) and an alias cannot exist with the same name as an index.

Table 3. Possible Values for `aliasReplacesIndex`

Value	Description
<code>true</code>	The index and it's content will be deleted and a versioned index and an alias will be created. Deleted content can be automatically reindexed on startup by using this in combination the <code>elasticSearch.bulkIndexOnStartup</code> config option
<code>false</code>	Falls back to the 'none' strategy. Event will be logged and a MappingException will be thrown.



The default is `true`.

## 2.4.3. `elasticSearch.migration.disableAliasChange`

In some cases the developer may prefer not to upgrade the alias to the new version of the index until some other tasks are performed. This allows them to disable automatically pointing the alias to a new version of the index when this is created. Aliases can be changed later on manually or programmatically using `elasticSearchAdminService`

Table 4. Possible Values for `disableAliasChange`

Value	Description
<code>false</code>	Standard behaviour
<code>true</code>	Prevents the aliases to be changed to point to a new index



The default is `false`.

## 2.5. Dynamic Method Injection

### 2.5.1. `elasticSearch.searchMethodName`

Change the name of search method in domain class. By default it's `search`.

For example

```
MyDomain.search("${params.query}")
```



In order to change the method name to `esSearch` just update the `elasticSearch.searchMethodName='esSearch'` in `application.groovy`

### 2.5.2. `elasticSearch.countHitsMethodName`

Change the name of `countHits` method in domain class. By default it's `countHits`.

For example

```
MyDomain.countHits("${params.query}")
```



In order to change the method name to `esCountHits` just update the `elasticSearch.countHitsMethodName='esCountHits'` in `application.groovy`

### 2.5.3. `elasticSearch.disableDynamicMethodsInjection`

To complete disabled injection of dynamic methods set `elasticSearch.disableDynamicMethodsInjection = true` in `applicaiton.groovy`

## 2.6. Others properties

### 2.6.1. `elasticSearch.datastoreImpl`

Only required when enabling the auto-index feature. This property specifies which GORM datastore implementation should be watched for storage events. The value should be the name of the datastore bean as it is configured in the Spring context; some possible values:

Table 5. Possible Values for `datastoreImpl`

Value	Description
<code>mongoDatastore</code>	The name of the MongoDB datastore bean.
<code>hibernateDatastore</code>	The name of the Hibernate datastore bean.

### 2.6.2. `elasticSearch.bootstrap.config.file`

When using then plugin to construct a local node, the default Elasticsearch configuration is used by default. If you use a modified Elasticsearch configuration, you can use this property to specify the location of the file (as an application resource).

### 2.6.3. `elasticSearch.bootstrap.transportSettings.file`

When choosing transport mode this configuration will be used to set up the `TransportClient` settings (used by some cloud providers).

#### 2.6.4. `elasticSearch.client.transport.sniff`

Only usable in with a transport client. Allows to sniff the rest of the cluster, and add those into its list of machines to use. In this case, the ip addresses used will be the ones that the other nodes were started with (the “publish” address)

#### 2.6.5. `elasticSearch.cluster.name`

The name of the cluster for the client to join.

#### 2.6.6. `elasticSearch.date.formats`

List of date formats used by the JSON unmarshaller to parse any date field properly. Note : future version of the plugin may change how formats are manipulated.

#### 2.6.7. `elasticSearch.defaultExcludedProperties`

List of domain class properties to automatically ignore (will not be indexed) if their name match one of those. This will apply to both the default-mapped domain class, with the static `searchable` property set to "true", and when using closure mapping. To override this setting on a specific class, it can be added to the `only` property of the `searchable` closure.

#### 2.6.8. `elasticSearch.disableAutoIndex`

A boolean determining if the plugin should reflect any database save/update/delete automatically on the indices. Default to `false`.

#### 2.6.9. `elasticSearch.bulkIndexOnStartup`

Determines whether the application should launch a bulk index operation upon startup.

Table 6. Possible Values for `bulkIndexOnStartup`

Value	Description
<code>false</code>	No indexing will happen on startup.
<code>true</code>	All content will be indexed on startup.
<code>'deleted'</code>	This value is related to the mapping migration strategy chosen. If any migration is required and any content is deleted due to it, on startup only indices and mappings lost will be indexed. More on <a href="#">Mapping Migrations</a> .



Default to `true`.

#### 2.6.10. `elasticSearch.index.name`

A string indicating which ElasticSearch index should be used. If not present, will default to the package name of the domain in question.



From Elasticsearch 5.0 on only selected settings like for instance `index.codec` can be set on the node level. All other settings must be set on each individual index. To set default values on every index, index templates should be used instead. So, unset this value if you are using Elasticsearch v5.0 or above.

### 2.6.11. `elasticsearch.index.compound_format`

Should the compound file format be used (boolean setting). Set to `false` by default (really applicable for file system based index storage). More details on this setting on the [ElasticSearch Documentation](#).

### 2.6.12. `elasticsearch.index.store.type`

Determine how the indices will be stored. More details on the possible values on the [ElasticSearch Documentation](#).

Table 7. Possible value for index store type

Value	Description
memory	Stores the index in memory. Useful for testing. <b>Please note</b> that it's not support in ES 2.0 and later versions.
mmapfs	Stores the shard index on the file system (maps to Lucene MMapDirectory) using mmap.
niofs	Stores the shard index on the file system (maps to Lucene NIOFSDirectory) and allows for multiple threads to read from the same file concurrently.
simplefs	Stores using a plain forward implementation of file system storage (maps to Lucene SimpleFsDirectory) using random access file.

### 2.6.13. `elasticsearch.index.settings.numberOfReplicas`

Sets the number of replicas created for each shard of the index. If not present, will default to zero.

### 2.6.14. `elasticsearch.gateway.type`

Determine the gateway type to be used. More details on the possible values are in the [ElasticSearch Documentation](#). Using a setting of "none" (possibly in combination with `index.store.type` set to "memory") can be useful for tests.

### 2.6.15. `elasticsearch.maxBulkRequest`

Max number of requests to process at once. Reduce this value if you have memory issue when indexing a big amount of data at once. If this setting is not specified, 500 will be use by default.

### 2.6.16. `elasticSearch.path.data`

The location of the data files of each index / shard allocated on the node.

### 2.6.17. `elasticSearch.path.plugins`

The location of plugin files such as native scripts. Each plugin will be contained in a subdirectory.

### 2.6.18. `elasticSearch.searchableProperty.name`

The name of the ElasticSearch mapping configuration property that annotates domain classes. The default is 'searchable'.

### 2.6.19. `elasticSearch.includeTransients`

Whether to index and search all non excluded transient properties. All explicitly included transients in `only` will be indexed regardless.



Default is `false`.

## Chapter 3. Quick Start

- Install Elasticsearch using instructions from [here](#).
- Just like any other Grails plugin, through the Grails Plugin center and Edit your project's build.gradle file, by adding the plugin's dependency declaration:

```
dependencies {  
    ...  
    compile "org.grails.plugins:elasticsearch:2.7.4"  
    ...  
}
```

- Overriding Spring Boot elasticsearch version

Spring Boot 1.3.x supports Elasticsearch 1.5.2 OOTB (<https://github.com/spring-projects/spring-boot/blob/master/spring-boot-dependencies/pom.xml#L76>) and will install dependencies for this version, if not explicitly overridden. To do so add the following to your **build.gradle**:

```
def elasticsearchVersion = '7.5.1'  
ext['elasticsearch.version'] = elasticsearchVersion
```

- Updated **application.yml** with default configurations:

```
elasticsearch:  
  datastoreImpl: hibernateDatastore  
  client:  
    mode: transport  
    hosts:  
      - {host: localhost, port: 9300}  
  cluster.name: <ENTER CLUSTER NAME HERE>
```

In order debug Elasticsearch plugin, add **logger("grails.plugins.elasticsearch", DEBUG, ['STDOUT'], false)** in **logback.groovy** file as following:

```

import grails.util.BuildSettings
import grails.util.Environment
import org.springframework.boot.logging.logback.ColorConverter
import org.springframework.boot.logging.logback.WhitespaceThrowableProxyConverter

import java.nio.charset.Charset

conversionRule 'clr', ColorConverter
conversionRule 'wex', WhitespaceThrowableProxyConverter

// See http://logback.qos.ch/manual/groovy.html for details on configuration
appender('STDOUT', ConsoleAppender) {
    encoder(PatternLayoutEncoder) {
        charset = Charset.forName('UTF-8')

        pattern =
            '%clr(%d{yyyy-MM-dd HH:mm:ss.SSS}){faint} ' + // Date
            '%clr(%5p) ' + // Log level
            '%clr(---){faint} %clr([%15.15t]){faint} ' + // Thread
            '%clr(%-40.40logger{39}){cyan} %clr(:){faint} ' + // Logger
            '%m%n%wex' // Message
    }
}

def targetDir = BuildSettings.TARGET_DIR
if (Environment.isDevelopmentMode() && targetDir != null) {
    appender("FULL_STACKTRACE", FileAppender) {
        file = "${targetDir}/stacktrace.log"
        append = true
        encoder(PatternLayoutEncoder) {
            pattern = "%level %logger - %msg%n"
        }
    }
    logger("StackTrace", ERROR, ['FULL_STACKTRACE'], false)
    logger("grails.plugins.elasticsearch", DEBUG, ['STDOUT'], false)
    root(ERROR, ['STDOUT', 'FULL_STACKTRACE'])
}
else {
    root(ERROR, ['STDOUT'])
}

```

### 3.1. Default mapping

To declare a domain class to be searchable, the simplest way is to define the following static property in the code:

```
static searchable = true
```



The plugin will generate a default mapping for each properties of the domain.

### 3.2. Custom mapping

You can customize how each properties are mapped to the index using a closure. The syntax is similar to GORM's mapping DSL.

```
static searchable = {  
    // mapping DSL...  
}
```

See below for more details on the mapping DSL.

### 3.3. Limit properties with only/except

**only** and **except** are used to limit the properties that are made searchable. You may not define both except & only settings at the same time.

The following code will only map the 'message' property, any others will be ignored.

```
class Tweet {  
    static searchable = {  
        only = 'message'  
    }  
    String message  
    String someUselessField  
}
```

The following code will map all properties except the one specified.

```
class Tweet {  
    static searchable = {  
        except = 'someUselessField'  
    }  
    String message  
    String someUselessField  
}
```

You can use a Collection to specify several properties.

```
class Tweet {
    static searchable = {
        except = ['someUselessField', 'userName']
    }
    String message
    String userName
    String someUselessField
}
```



The properties that are ignored will not be sent to Elasticsearch. It also means that when you will get back a domain from Elasticsearch, some fields that are not supposed to be null, may still be null.

### 3.4. Including transients

How the plugin manages transient properties is controlled by the `elasticsearch.includeTransients` configuration property. If this is set to `false` only transient properties explicitly included in `only` will be mapped and searchable, if set to `true`, all domain class properties will be mapped, including `transients`.

The following are valid examples

```
//assert grailsApplication.config.elasticsearch.includeTransients == false
class Person {
    String firstName
    String lastName
    String getFullName() {
        firstName + " " + lastName
    }
    static transients = ['fullName']
    static searchable = {
        only = ['fullName']
    }
}

// new Person(firstName: "Nikola", lastName: "Tesla")
// can be found using:
// def tesla = Person.search("Nikola Tesla").searchResults.first()
```

```
//assert grailsApplication.config.elasticSearch.includeTransients == true
class Multiplication {
    int opA
    int opB
    int getResult() {
        opA * opB
    }
    static transients = ['result']
    static searchable = true
}
// new Multiplication(opA: 2, opB: 3)
// can be found using:
// def multiplication = Multiplication.search("2").searchResults.first()
// def multiplication = Multiplication.search("3").searchResults.first()
// def multiplication = Multiplication.search("6").searchResults.first()
```



From the examples above, once the domain object is found, its transient values will be calculated from the information stored on ElasticSearch: `multiplication.result == 6`, but `tesla.fullName == "null null"`, as `firstName` and `lastName` were not indexed. This behaviour can be prevented by creating convenient setters for the transient properties.

### 3.5. Transients and collections

When transient properties are collections the only way the plugin can define the correct ElasticSearch mapping during boot is if the element types are explicitly defined on the grails domain object. For instances of `Collection` this can be achieved by defining its type on the `hasMany` property (otherwise the ElasticSearch type will be defined as `object`). This is not required for arrays.

Some valid examples:

```
class Tweet {
    String message
    List getHashtags() { ... }
    static transients = ['hashtags']
    static hasMany = [hashtags: String]
    static searchable = {only = 'hashtags' }
}
```

```
class FamilyGuy {  
    String wife  
    String son  
    String daughter  
    String baby  
    String[] getRelatives() { ... }  
    static transients = ['relatives']  
    static searchable = { only = 'relatives' }  
}
```

### 3.6. The Mapper Attachment plugin

The Elasticsearch Mapper Attachment plugin is deprecated since ES version 5.0.0. Hence it is no longer supported in this plugin.

In future versions it could be possible to support the new [Ingest Attachment Processor Plugin](#).

# Chapter 4. Mapping

From version 0.0.4.0 in addition to the indices generated by the plugin based on the domain objects package names or configuration, two new aliases are created for every index: `<indexName>_read` and `<indexName>_write`. These two aliases are being used by the plugin to index and query from Elasticsearch and are needed to centralise the choice of index to use during mapping migrations when the `'alias'` strategy is being used and there are multiple instances of the application.

## 4.1. Class Mapping

### 4.1.1. root

Determine if the domain class will have its own index or not. Take a boolean as parameter, and is set to `true` by default.

```
class Preference {
    static searchable = {
        root false
    }
    // ...
}

class Tag {
    static searchable = true
    // ...
}

class Tweet {
    static searchable = {
        message boost:2.0
    }
    // ...
}
```

In this code, the classes `Tweet` and `Tag` are going to have their own index. The class `Preference` will not. It also means that any search request will never return a `Preference`-type hit. The dynamic method `search` will not be injected in the `Preference` domain class.

The domains not root-mapped can still be considered searchable, as they can be components of another domain which is root-mapped. For example, considered the following domain:

```
class User {
  static searchable = {
    userPreferences component:true
  }

  Preference userPreferences
}
```

When searching, any matches in the `userPreferences` property will be considered as a `User` match.

### 4.1.2. all

Set default analyzer for all domain class fields.

```
static searchable = {
  all = [analyzer: 'russian_morphology']
}
```

```
static searchable = {
  all = false
}
```

When disabling the all field, it is a good practice to set `index.query.default_field` to a different value (for example, if you have a main 'message' field in your data, set it to message).

## 4.2. Properties Mapping

You can customize the mapping for each domain properties using the closure mapping. The syntax is simple:

```
static searchable = {
  propertyName option1:value, option2:value, ...
}
```

### 4.2.1. Available options

Option Name	Values	Description
boost	Number	A decimal boost value. With a positive value, promotes search results for hits in this property; with a negative value, demotes search results that hit this property.

Option Name	Values	Description
<code>component</code>	<code>true, false</code>	To use only on domain (or collection of domains), make the property a searchable component.
<code>converter</code>	A <code>Class</code>	A <code>Class</code> to use as a converter during the marshalling/un-marshalling process for that peculiar property. That class must extends the <code>PropertyEditorSupport</code> java class.
<code>excludeFromAll</code>	<code>true, false</code>	determines if the property is to append in the " <code>_all</code> " field. Default to <code>true</code> .
<code>index</code>	<code>false, true.</code>	The index option controls whether field values are indexed. It accepts <code>true</code> or <code>false</code> and defaults to <code>true</code> . Fields that are not indexed are not queryable.
<code>fielddata</code>	<code>true, false</code>	Enables the use of text fields in sorting and aggregation. Be careful as fielddata can consume a lot of heap space, especially when loading high cardinality text fields. Default <code>false</code>
<code>reference</code>	<code>true, false</code>	To use only on domain (or collection of domains), make the property a searchable reference.
<code>parent</code>	<code>true, false</code>	A boolean value to be used in conjunction with the <code>reference</code> or <code>component</code> property . Set to <code>true</code> if the referenced field should be mapped as the parent of this document. Default set to <code>false</code> .
<code>multi_field</code>	<code>true, false</code>	A boolean value. Maps the value of the field twice; Once with it being analyzed, and once with it being not_analyzed under the name untouched. Default set to <code>false</code> .

Option Name	Values	Description
fields	A map describing the extra fields	Maps the value of the field more than once; it is useful when you want to index a field with different analyzers. See the <a href="#">ElasticSearch documentation</a> . Default set to <b>null</b> .
<a href="#">geoPoint</a>	<b>true, false</b>	Maps the field to a geo_point. Default: <b>false</b>
<a href="#">alias</a>	String	A string value. The field noted with this parameter will be duplicated to an alias
<a href="#">dynamic</a>	<b>true, false</b>	Only available for String properties. Determines whether this field should be dynamically mapped by elasticsearch.

## 4.3. Parent Child

To map a parent/child relationship, the child element must either contain the parent element as a component or reference it as a referenced document. This component must be mapped as a parent in the child element.

### Example

```
class ParentElement {
  ...
}

class EmbeddingChild {
  ParentElement parentElement

  static searchable = {
    parentElement parent: true, component: true
  }
}

class ReferencingChild {
  ParentElement parentElement

  static searchable = {
    parentElement parent: true, reference: true
  }
}
```



## 4.4. Geo Point

A geographic location can be mapped to a `geo_point`. The field for the longitude has to be named `lon` and the field for the latitude has to be named `lat`

### Example

```
class GeoPoint {  
  
    Double lat  
    Double lon  
  
    static searchable = {  
        root false  
    }  
}  
  
class Building {  
  
    String name  
    GeoPoint location  
  
    static searchable = {  
        location geoPoint: true, component: true  
    }  
}
```

## 4.5. Alias

A field can be aliased. This is useful in situations where another service may expect certain tags.

For example, [Kibana](#) uses an `\@timestamp` field to filter report records by date.

### Example

```
class Session {  
  
    Date loginTime  
  
    static searchable = {  
        loginTime alias:'@timestamp'  
    }  
}
```

## 4.6. Dynamic

Elasticsearch can map field contents as dynamic objects.

This is especially useful if you store JSON Strings in your database and want to make those objects searchable in elasticsearch.

### Example

```
class Session {  
  
    String jsonData  
  
    static searchable = {  
        dynamic: true  
    }  
}  
  
Session session = new Session()  
session.jsonData = ([foo: 'bar'] as JSON).toString()
```

The default mapping would make the jsonData field an escaped String field and a search for jsonData.foo = bar would result in no result. With dynamic mapping enabled for this field, we enable JSON handling of this field and tell elasticsearch to map this field dynamically. The result is that a search for `jsonData.foo=bar` would result in a search hit.



This will only work on String fields and will result in an error if the String is no valid json

## 4.7. Searchable Component Reference

The plugin support a similar searchable-component & searchable-reference behaviour from Compass when you are dealing with domain association. See below to find out about the difference between both mapping modes.

### 4.7.1. Searchable Reference

The searchable-reference mapping mode is the default mode used for association, and requires the searchable class of the association to be root-mapped in order to have its own index. With this mode, the associated domains are not completely marshalled in the resulting JSON document: only the id and the type of the instances are kept. When the document is retrieved from the index, the plugin will automatically rebuild the association from the indices using the stored id.

### Example

```

class MyDomain {
    // odom is an association with the OtherDomain class, set as a reference
    OtherDomain odom

    static searchable = {
        odom reference:true
    }
}

// The OtherDomain definition, with default searchable configuration
class OtherDomain {
    static searchable = true

    String field1 = "val1"
    String field2 = "val2"
    String field3 = "val3"
    String field4 = "val4"
}

```

When indexing an instance of `MyDomain`, the resulting JSON documents will be sent to ElasticSearch:

```

{
  "mydomain": {
    "_id":1,
    "odom": { "id":1 }
  }
}

{
  "otherdomain": {
    "_id":1,
    "field1":"val1",
    "field2":"val2",
    "field3":"val3",
    "field4":"val4"
  }
}

```

### 4.7.2. Searchable Component

The searchable-component mapping mode must be explicitly set, and does not require the searchable class of the association to be root-mapped.

With this mode, the associated domains are nested in the parent document.

#### Example

```

class MyDomain {
    // odom is an association with the OtherDomain class, set as a reference
    OtherDomain odom

    static searchable = {
        odom component:true
    }
}

// The OtherDomain definition, with default searchable configuration
class OtherDomain {
    static searchable = true

    String field1 = "val1"
    String field2 = "val2"
    String field3 = "val3"
    String field4 = "val4"
}

```

When indexing an instance of MyDomain, the resulting JSON document will be sent to ElasticSearch:

```

{
  "mydomain": {
    "_id":1,
    "odom": {
      "_id":1,
      "field1":"val1",
      "field2":"val2",
      "field3":"val3",
      "field4":"val4"
    }
  }
}

```

If you'd rather that the reference object be mapped with type 'inner' rather than the default 'nested', set the 'component' key with a value of 'inner' rather than 'true':

```

class MyDomain {
    // odom is an association with the OtherDomain class, set as a reference
    OtherDomain odom

    static searchable = {
        odom component: 'inner'
    }
}

```

## 4.8. Mapping Migrations

During the application startup the application will attempt to create the needed indices on Elasticsearch and create the type mappings defined by the user. If these indices and mappings already existed on the Elasticsearch cluster (ie. an older version of the application was running against it) and the new mapping definitions differ with the existing ones there's the potential for a Mapping conflict. This section describes how to configure the application to deal with this scenario.

It is important to highlight that not all type mapping changes will result on a conflict. Ie. adding a new field to a mapping does not result in a conflict whilst changing a property from component:'inner' to nested or vice-versa, will. These strategies will only be needed and applied when a **conflicting** mapping is being found.

### 4.8.1. Migration Strategies

The migration strategy is being defined by the `elasticSearch.migration.strategy` configuration property, and it accepts three values:

- `'none'`
- `'deleteIndex'`
- ~~`'delete'` is no longer supported since the upgrade to Elasticsearch > 2.0~~
- `'alias'`

The default strategy is `'alias'` as it is the only strategy that can achieve zero-downtime migrations and thus [recommended by Elasticsearch](#)

These values are described on more detail further ahead

### 4.8.2. Migration Strategy 'none'

This option keeps the original behaviour the plugin used before the Migration Strategies were implemented. When a Mapping Merge conflict is identified the event will be logged and an Exception will be logged. It will be responsibility for the application administrator to manually fix the problem.

This configuration was left as a backwards compatibility and it will prevent the application from booting successfully, therefore we **discourage teams from using this option**.

### 4.8.3. Migration Strategy 'deleteIndex'

When choosing this option, when a conflict occurs installing mapping, the application will delete the existing index for the type, **alongside with all content indexed on that index (including content from other mappings)** and it will recreate the index and all its mappings. There are a couple of important details on this information:

- Deleted documents can be automatically reindexed on startup by using the `elasticSearch.bulkIndexOnStartup` configuration property (See below)
- New indices will be created with a version number and the right aliases, to make them

compatible with potential future 'alias' migrations (without requiring additional index deletions)

- Using this configuration there will always be a time window (between deletion and reindexation) where documents can't be found by search, therefore this option cannot achieve a **zero-downtime** deployment

See [Dealing with deleted content](#) below for more details on automatic indexing.

#### 4.8.4. Migration Strategy 'alias'

This is the migration strategy [recommended by Elasticsearch](#).

To better understand this strategy we will describe a typical 'alias' migration.

Elasticsearch contains

```
index 'myapplication.store_v27' with types 'car' and 'motorbike'
alias 'myapplication.store' pointing to 'myapplication.store_v27'
'myapplication.store_v27/car' contains 520 documents
'myapplication.store_v27/motorbike' contains 12 documents
index 'myapplication.admin_v0' with type 'quote'
alias 'myapplication.admin' pointing to 'myapplication.admin_v0'
'myapplication.admin_v0/quote' contains 3200 documents
```

The application is configured to use indexes based on package names

'myapplication.store' and 'myapplication.admin'

(which as we already explained are actually aliases that point to versioned indices)

The team introduced a change on the Car domain that results in a conflict on the 'car' mapping

The application starts up

    Tries to install the mapping for 'motorbike', it detects the conflict

    Creates a new index called 'myapplication.store\_v28'

    Creates mappings 'myapplication.store\_v28/car' and

'myapplication.store\_v28/motorbike'

    Points all indexing requests for Car and Motorbike to the new index, while queries still happen on 'myapplication.store'

On Bootstrap (bulkIndexOnStartup)

    It indexes 520 cars into 'myapplication.store\_v28/car'

    It indexes 12 motorbikes into 'myapplication.store\_v28/motorbike'

    Switches the 'myapplication.store' alias to point to 'myapplication.store\_v28'

    Now all cars are indexed according to the new mapping

    Now all motorbikes are indexed according to the new mapping



All content can be queried at all times, during Bootstrap `bulkIndexOnStartup` content will be retrieved from the old index.



Even though there wasn't a conflict on 'car', all cars needed to be reindexed as they lived on the same index.

There are three potential scenarios when using the '[alias](#)' strategy:

Scenario	Behaviour
The index (ie. 'myapplication.store') does not exist	On this case there is not possibility of conflicts, as no previous mapping exist. However the application will behave slightly different than on the other two scenarios. Instead of creating the index (ie. 'myapplication.store'), it will create version 0 of it (ie. 'myapplication.store_v0') and an alias pointing to it. This is to facilitate the creation of future versions in case of conflict.
Alias exists pointing to a version (ie. 'myapplication.store' → 'myapplication.store_v27')	If there's a conflict on a mapping on the index, it will create a new version (ie. 'myapplication.store_v28'), reindex the content or not depending on the value of the <a href="#">elasticSearch.bulkIndexOnStartup</a> configuration property and point the alias to the new version once done.
Index already exists (ie. 'myapplication.store')	Elasticsearch cannot rename an index or create an alias with the same name as an index. The two alternatives here are to delete the index or fail the migration. This is controlled by the <a href="#">elasticSearch.migration.aliasReplacesIndex</a> configuration property, if set to true, it will delete the index and proceed the same way as when the index did not exist. The deleted documents will be reindexed or not depending on the value of the <a href="#">elasticSearch.bulkIndexOnStartup</a> . <b>This is the only scenario where there is content loss/downtime using the '<a href="#">alias</a>' strategy.</b>

In the case you wanted to create a new version of an index, but not change where the alias points to (ie. for testing or if you wanted to perform extra tasks on the index before updating the alias), the [elasticSearch.migration.disableAliasChange](#) configuration property can be used



Aliases will only point to the new version of the index once all content is reindexed (if chosen to). Meanwhile, all index requests, either by [elasticSearchService](#) or using dynamic finders will go to the new version of the index, whilst queries will go to the old version of the index.

See [Dealing with deleted content](#) below for more details on automatic indexing.

#### 4.8.5. Dealing with deleted content

Using the '[delete](#)' or '[alias](#)' strategy may lead to deleting content stored on Elasticsearch. This

content can be automatically reindexed using the `elasticSearch.bulkIndexOnStartup`. The duration of this process will depend on the amount of content to index.

When this property is set to `true` all content will be deleted. When set to `'deleted'` only the domain classes which documents were deleted will be indexed. In either case, when using the `'alias'` strategy, once all content is indexed all aliases will point to the latest version of the index.



# Chapter 5. Indexing

With its default configuration (with the `disableAutoIndex` configuration key set to `false`), the plugin is indexing automatically any searchable domains when GORM/Hibernate do a save or an update in the database.

It also deletes automatically from the index any document corresponding to a domain that is deleted from the database. You normally shouldn't have to worry about indexing, but sometimes you may have to do it by yourself, for example on dirty domain object that you may not want to save right now.

The plugin is providing a few injected methods in the domain or in the `ElasticSearchService` to allow that.

## 5.1. Index examples

```
// Index all searchable instances
elasticSearchService.index()

// Index a specific domain instance
MyDomain md = new MyDomain(value:'that')
md.save()
elasticSearchService.index(md)

// Index a collection of domain instances
def ds = [new MyDomain(value:'that'), new MyOtherDomain(name:'this'), new
MyDomain(value:'thatagain')]
ds*.save()
elasticSearchService.index(ds)

// Index all instances of the specified domain class
elasticSearchService.index(MyDomain)
elasticSearchService.index(class:MyDomain)
elasticSearchService.index(MyDomain, MyOtherDomain)
elasticSearchService.index([MyDomain, MyOtherDomain])
```

## 5.2. Unindex examples

```
// Unindex all searchable instances
elasticSearchService.unindex()

// Unindex a specific domain instance
MyDomain md = new MyDomain(value:'that')
md.save()
elasticSearchService.unindex(md)

// Unindex a collection of domain instances
def ds = [new MyDomain(value:'that'), new MyOtherDomain(name:'this'), new
MyDomain(value:'thatagain')]
ds*.save()
elasticSearchService.unindex(ds)

// Unindex all instances of the specified domain class
elasticSearchService.unindex(MyDomain)
elasticSearchService.unindex(class:MyDomain)
elasticSearchService.unindex(MyDomain, MyOtherDomain)
elasticSearchService.unindex([MyDomain, MyOtherDomain])
```

# Chapter 6. Searching

The plugin provides 2 ways to send search requests.

- You can use the `elasticSearchService` and its public `search` method for cross-domain searching, meaning that ElasticSearch may analyze multiple indices and return hits of different types (=different domains).

```
def res = elasticSearchService.search("${params.query}")
// 'res' search results may contains multiple types of results
```

- You can use the injected dynamic method in the domain for domain-specific searching.

```
def res = Tweet.search("${params.query}")
// 'res' search results contains only Tweet instances
```

These search methods return a `Map` containing 3 entries:

- a `total` entry, representing the total number of hits found
- a `searchResults` entry, containing the hits
- a `scores` entry, containing the hits scores

## Example

```
def res = Tweet.search("${params.query}")
println "Found ${res.total} result(s)"
res.searchResults.each {
    println it.message
}

def res = elasticSearchService.search("${params.query}")
println "Found ${res.total} result(s)"
res.searchResults.each {
    if(it instanceof Tweet) {
        println it.message
    } else {
        println it.toString()
    }
}
```

If you're willing to retrieve only the number of hits for a peculiar query, you can use the `countHits()` method. It will only return an `Integer` representing the total hits matching your query.

## Example

```
def res = Tweet.countHits("${params.query}")
println "Found ${res} result(s)"

def res = elasticSearchService.countHits("${params.query}", [indices:'test'])
println "Found ${res} result(s)"
```

## 6.1. Query Strings

The search method injected in the domain or the `ElasticSearchService` has multiple signatures available. You can pass it a simple `String` to compute your search request. That string will be parsed by the **Lucene query parser** so feel free to use its syntax to do more specific search query.

You can find out about the syntax on the [Apache Lucene website](#).

### Example

```
def results = elasticSearchService.search("${params.query}")
def resultsTweets = Tweet.search("message:${params.query}")
```

## 6.2. Query Closure

You can use the [Groovy Query DSL](#) to build your search query as a `Closure`.

The format of the search `Closure` follow the same JSON syntax as the [ElasticSearch REST API](#) and the [Java Query DSL](#).

### Example

```
def result = elasticSearchService.search(searchType: 'dfs_query_and_fetch') {
  bool {
    must {
      query_string(query: params.query)
    }
    if (params.firstname) {
      must {
        term(firstname: params.firstname)
      }
    }
  }
}
```

## 6.3. Query Builder

A `QueryBuilder` can be passed to the search method.

## Example

```
QueryBuilder query = QueryBuilders.matchAllQuery()
def result = elasticSearchService.search(query)
```

## 6.4. Filter Closure

A filter closure can be passed as a second argument after the search closure to the search method.

## Example

```
def result = elasticSearchService.search(
    [indices: Building, types: Building, sort: sortBuilder],
    null as Closure,
    {
        geo_distance(
            'distance': '5km',
            'location': [lat: 48.141, lon: 11.57]
        )
    })
```

## 6.5. Filter Builder

A **FilterBuilder** filter can be passed as a second argument after the search parameter to the search method.

## Example

```
FilterBuilder filter = FilterBuilders.rangeFilter("price").gte(1.99).lte(2.3)
def result = elasticSearchService.search(
    [indices: Building, types: Building, sort: sortBuilder],
    null as Closure,
    filter)
```

## 6.6. Highlighting

The search method support highlighting: automatic wrapping of the matching terms in the search results with HTML/XML/Whatever tags.

You can activate this with a **Closure** containing the highlight settings in the search method **highlight** parameter.

The format of the **Closure** for defining the highlight settings is the same as the [ElasticSearch REST API](#).

## Example

```
// Define the pre & post tag that will wrap each term matched in the document.
def highlighter = {
  field 'message'
  field 'tags.name'
  preTags '<strong>'
  postTags '</strong>'
}

def results = Tweet.search("${params.query}", [highlight: highlightSettings])
```

### 6.6.1. Highlight results

If a search result is found, the `search` method will add a `highlight` entry in the map result.

That entry contains a `List` with every highlighted fragments/fields found for each hit.

```
def results = Tweet.search("${params.query}", [highlight: { field 'message' }])
def highlighted = results.highlight

results?.searchResults?.eachWithIndex { hit, index ->
  // Retrieve the 'message' field fragments for the current hits
  def fragments = highlighted[index].message?.fragments

  // Print the fragment
  println fragments?.size() ? fragments[0] : ''
}
```

### 6.6.2. Highlighted fields

To determine which fields are to be processed by ElasticSearch, use the `field` setting.

You can call the `field` setting as many time as you want to add any field.

#### Signature

```
field <fieldName>, <fragmentSize>, <numberOfFragment>
```

#### Examples

```
def highlightSettings = {
    field 'message' // Add the 'message' field in the highlighted
    fields list
    field 'tags.name' // Add the 'name' field contained in the 'tags'
    field of // the document in the highlighted fields list
    field 'thatAwesomeField', 0, 20 // Add the 'thatAwesomeField' field with
    // some values fixed for fragmentSize and
    // numberOfFragment parameters
}

def highlightSettings2 = {
    field '_all' // Add the special '_all' field in the
    highlighted // fields list
}

def results = Tweet.search("${params.query}", [highlight: highlightSettings])
def results2 = Tweet.search("${params.query}", [highlight: highlightSettings2])
```

### 6.6.3. Highlighting tags

By default, Elasticsearch will use emphasis tag “<em>...</em>” to wrap the matching text.

You can customize the tags with the `preTags` and `postTags` settings.

```
def highlightSettings = {
    field 'message'
    preTags '<myAwesomeTag>'
    postTags '</myAwesomeTag>'
}
```

## 6.7. Sorting

To sort the search results, either a field name or a `SortBuilder` must be passed.

### Returned sort values

The sort values are not part of the search results themselves but are part of `result.sort`.

`sort` contains all search values calculated by the Elasticsearch server as a list mapped to the id of the respective domain objects

### Example

```
assert [1:[23, 42], 2: [24, 40]] == result.sort
```

Unresolved directive in searching/sorting.adoc - include::sorting/geoDistanceSorting.adoc[]

## 6.8. Pagination

To return paginated results one can pass **from** and **size** to search method query params.

### Example

```
Post.search('xyz', [ indices: Post, types: Post, from: 0, size: 10 ] )

// or

elasticSearchService.search('xyz', [ indices: Post, types: Post, from: 0, size: 10 ])
```

## 6.9. Aggregation

You can use AggregationBuilders while searching. The aggregations' framework helps provide aggregated data based on a search query. It is based on simple building blocks called aggregations, that can be composed in order to build complex summaries of the data.

You could use preconstructed AggregationBuilders or use the standard DSL based aggregations

### Examples

*Use of AggregationBuilders*



```

// build your aggregation(s)
def aggregations = []
aggregations << AggregationBuilders.filters('types',
    new FiltersAggregator.KeyedFilter('jpg', QueryBuilders.matchQuery('type',
'jpg'))),
    new FiltersAggregator.KeyedFilter('png', QueryBuilders.matchQuery('type',
'png'))
)
aggregations << AggregationBuilders.terms('names').field('name')
aggregations << AggregationBuilders.avg('avg_size').field('size')

// Use it

Photo.search('xyz', null as Closure, aggregations)

// or

elasticSearchService.search(
    'xyz',
    null as Closure,
    aggregations,
    [ indices: Photo, types: Photo ])

```

*Use of DSL*

```

Photo.search('xyz', null as Closure, {
    "types" {
        filters {
            "filters" {
                "jpg" { match(type: 'jpg') }
                "png" { match(type: 'png') }
            }
        }
    }

    "names" {
        terms(field: 'name')
    }

    "avg_size" {
        avg(field: 'size')
    }
})

```

// or

```

elasticSearchService.search(
    'xyz',
    null as Closure,
    {
        "types" {
            filters {
                "filters" {
                    "jpg" { match(type: 'jpg') }
                    "png" { match(type: 'png') }
                }
            }
        }

        "names" {
            terms(field: 'name')
        }

        "avg_size" {
            avg(field: 'size')
        }
    },
    [ indices: Photo, types: Photo ])

```

# Chapter 7. Admin

The plugin implements a few convenience methods for a few admin-oriented actions.

## 7.1. Refresh

Explicitly refresh one or more index, making all operations performed since the last refresh available for search. It will also flush the current `IndexRequestQueue` if there are pending index or delete requests from the application side.

The refresh method is not asynchronous, meaning that it will wait for all operations to complete before resuming the execution of your application.

```
elasticSearchService.index(domain)
// Some code...
// ...

elasticSearchService.index(domain2)
// Some code...
// ...

elasticSearchService.index(domain3)
// Some code...
// ...

elasticSearchAdminService.refresh() // Ensure that the 3 previous index
                                   // requests have been made searchable by ES
```

## 7.2. Delete Index

Delete an index, all its mapping and its content from the `ElasticSearch` instance. Be careful when using this command because it cannot be undone.



The generated mapping from the grails plugin is also deleted.

The method can be limited to one or more specific indices or applied to all indices at once (called with no parameter).

```
elasticSearchAdminService.deleteIndex()
```

## Chapter 8. Low Level API

If you need to use the Elastic Search client directly, you can use the `elasticSearchHelper` bean that is injected in any services/controllers to get the current instance.

Simply encapsulate your code within a `withElasticSearch` bloc, and you will get a `org.elasticsearch.client.Client` implementation to play with.

```
class MySearchService {
  static transactional = true

  def elasticSearchHelper

  def myMethod(indexName, settings) {
    elasticSearchHelper.withElasticSearch { client ->
      // Do some stuff with the ElasticSearch client
      client.admin()
        .indices()
        .prepareCreate(indexName)
        .setSettings(settings)
        .execute()
        .actionGet()
    }
  }
}
```

Please refer to the [Elastic Search API](#) for more information on the methods and properties available on the client.

# Chapter 9. Example

## 9.1. Twitter

### 9.1.1. The Domains

```
class Tweet {  
  static searchable = {  
    message boost:2.0  
  }  
  
  static belongsTo = [  
    user:User  
  ]  
  
  static hasMany = [  
    tags:Tag  
  ]  
  
  static constraints = {  
    tags nullable:true, cascade:'save, update'  
  }  
  
  String message = ''  
  Date dateCreated = new Date()  
}
```

```

class User {
    static searchable = {
        except = 'password'
        lastname boost:20
        firstname boost:15, index:'true'
        listOfThings index:'false'
        someThings index:'false'
        tweets component:true
    }

    static constraints = {
        tweets cascade:'all'
    }
    static hasMany = [
        tweets:Tweet
    ]
    static mappedBy = [
        tweets:'user'
    ]

    String lastname
    String firstname
    String password
    String activity = 'Evildoer'
    String someThings = 'something'
    ArrayList<String> listOfThings = ['this', 'that', 'andthis']
}

```

```

class Tag {
    static searchable = {
        except=['boostValue']
    }

    String name
    Integer boostValue = 1
}

```

### 9.1.2. The Controller

- A action triggering indexation

`ElasticSearchController` (`testCaseService` is just dealing with GORM instructions):

```

class ElasticSearchController {
  def elasticSearchService
  def testCaseService

  def postTweet = {
    if(!params.user?.id) {
      flash.notice = "No user selected."
      redirect(action: 'index')
      return
    }
    User u = User.get(params.user.id)
    if (!u) {
      flash.notice = "User not found"
      redirect(action: 'index')
      return
    }
    // Create tweet
    testCaseService.addTweet(params.tweet?.message, u, params.tags)

    flash.notice = "Tweet posted"
    redirect(action: 'index')
  }
}

```

With this code (considering that there are already **User** in the database), new Tweets will be indexed automatically, and corresponding **User** indexed documents will be updated since we have set the **tweets** association as component.

- **Searching for Tweets**

```

def searchForUserTweets = {
  def tweets = Tweet.search("${params.message.search}").searchResults
  def tweetsMsg = 'Messages : '
  tweets.each {
    tweetsMsg += "<br />Tweet from ${it.user?.firstname} ${it.user?.lastname} :  

    ${it.message} "
    tweetsMsg += "(tags : ${it.tags?.collect{t -> t.name}})"
  }
  flash.notice = tweetsMsg
  redirect(action: 'index')
}

```

- **Searching for anything**

```

def searchAll = {
  def res = elasticSearchService.search("${params.query}").searchResults
  def resMsg = '<strong>Global search result(s):</strong><br />'
  res.each {
    switch(it){
      case Tag:
        resMsg += "<strong>Tag</strong> ${it.name}<br />"
        break
      case Tweet:
        resMsg += "<strong>Tweet</strong> \"${it.message}\" from  

${it.user.firstname} ${it.user.lastname}<br />"
        break
      case User:
        resMsg += "<strong>User</strong> ${it.firstname} ${it.lastname}<br />"
        break
      default:
        resMsg += "<strong>Other</strong> ${it}<br />"
        break
    }
  }
  flash.notice = resMsg
  redirect(action:'index')
}

```

## 9.2. Geo Distance Search

A search for buildings with a `geo_distance` filter, ordered by distance.

### 9.2.1. Domains

```

class GeoPoint {

  Double lat
  Double lon

  static searchable = {
    root false
  }
}

```

`GeoPoint` represents the geo coordinates for a building. The field names `lat` and `lon` are mandatory.



```
class Building {

    String name
    GeoPoint location

    static searchable = {
        location geoPoint: true, component: true
    }
}
```

The location of the building is mapped to an Elasticsearch [geo\\_point](#).

### 9.2.2. Service Methods

Searching for all buildings sorted by distance with 5km radius around geo location (lat=41.141, lon=11.57)

```
def searchForBuildings() {
    Closure filter = {
        geo_distance(
            'distance': '5km',
            'location': [lat: 48.141, lon: 11.57]
        )
    }

    def sortBuilder = SortBuilders.geoDistanceSort("location").
        point(48.141, 11.57).
        unit(DistanceUnit.KILOMETERS).
        order(SortOrder.ASC)

    def result = elasticSearchService.search(
        [indices: Building, types: Building, sort: sortBuilder],
        null as Closure,
        filter)

    return [results: result.searchResults, distances: result.sort]
}
```

The calculated distances are not part of the search results themselves but are part of **result.sort**. **sort** contains all search values calculated by the Elasticsearch server as a list mapped to the id of the respective domain objects

#### Example

```
assert [1:[23, 42], 2: [24, 40]] == result.sort
```

## 9.3. Parent/Child mapping

Mapping types being removed, please see: <https://www.elastic.co/guide/en/elasticsearch/reference/current/removal-of-types.html>

A store with many departments

```
class Store {  
  
    String name  
    String description = "A description of a store"  
    String owner = "Shopowner"  
  
    static searchable = true  
  
    static constraints = {  
        name blank: false  
        description nullable: true  
        owner nullable: false  
    }  
}  
  
class Department {  
  
    String name  
    Long numberOfProducts  
    Store store  
  
    static constraints = {  
        numberOfProducts nullable: true  
    }  
  
    static searchable = {  
        store join: true, component:true  
    }  
}
```

Search for all departments which are childs of a store with the owner "Shopowner"

```
def result = elasticSearchService.search(  
    QueryBuilders.hasParentQuery("store", QueryBuilders.matchQuery("owner",  
"Shopowner")),  
    null as Closure,  
    [indices: Department, types: Department]  
)
```

# Chapter 10. Frequently Asked Questions

## 10.1. Conflicting module versions

Normally this issues happen where different versions of same modules are download from different plugins or dependencies in `build.gradle`

It seems like there are multiple versions of groovy in the classpath. Looks similar to <http://stackoverflow.com/questions/31099214/tomcat-conflicting-module-versions-module-groovy-all-is-loaded-in-version-2-3>

### 10.1.1. Exception

```
log4j:ERROR Error initializing log4j: null
java.lang.ExceptionInInitializerError
    at org.codehaus.groovy.runtime.InvokerHelper.<clinit>(InvokerHelper.java:61)
    at groovy.lang.GroovyObjectSupport.<init>(GroovyObjectSupport.java:32)
    at org.codehaus.groovy.grails.commons.AbstractGrailsApplication.<init>(
        AbstractGrailsApplication.java:45)
    at org.codehaus.groovy.grails.commons.DefaultGrailsApplication.<init>
        (DefaultGrailsApplication.java:95)
    at org.codehaus.groovy.grails.commons.DefaultGrailsApplication.<init>
        (DefaultGrailsApplication.java:91)
    at sun.reflect.NativeConstructorAccessorImpl.newInstance0(Native Method)
    at sun.reflect.NativeConstructorAccessorImpl.newInstance(
        NativeConstructorAccessorImpl.java:62)
    at sun.reflect.DelegatingConstructorAccessorImpl.newInstance(
        DelegatingConstructorAccessorImpl.java:45)
    at java.lang.reflect.Constructor.newInstance(Constructor.java:423)
    at java.lang.Class.newInstance(Class.java:442)

    at
org.apache.tomcat.util.net.JIoEndpoint$SocketProcessor.run(JIoEndpoint.java:314)
    at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1142)
    at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:617)
    at
org.apache.tomcat.util.threads.TaskThread$WrappingRunnable.run(TaskThread.java:61)
    at java.lang.Thread.run(Thread.java:745)
**Caused by: groovy.lang.GroovyRuntimeException: Conflicting module versions. Module
[groovy-all is loaded in version 2.3.11 and you are trying to load version 2.4.4**
```

### 10.1.2. Workaround

```
grails.project.dependency.resolution = {
  plugins {
    compile ":elasticsearch:0.0.4.5", { exclude "groovy-all" } // elasticsearch
1.6.0
  }
}
```

- Above problem was reported with plugin version 0.0.4.5 and Grails 2.5.5.